

# From Pixels to Perception: Deep Learning with PyTorch

William Edward Hahn

September 5, 2023



# Contents

<b>1</b>	<b>Introduction to Data Science with PyTorch</b>	<b>1</b>
1.1	The Power of Data Science . . . . .	1
1.2	Tools for the Modern Data Scientist . . . . .	1
1.3	A Glimpse of Python Code . . . . .	1
1.4	Structure of This Book . . . . .	2
1.5	Recap . . . . .	2
<b>I</b>	<b>Fundamentals of Data Science</b>	<b>3</b>
<b>2</b>	<b>Understanding Data</b>	<b>5</b>
2.1	What is Data? . . . . .	5
2.2	Types of Data . . . . .	5
2.3	Data Cleaning . . . . .	5
2.4	Data Visualization . . . . .	6
2.5	Recap . . . . .	6
<b>II</b>	<b>NumPy for Numerical Analysis</b>	<b>9</b>
<b>3</b>	<b>Getting Started with NumPy</b>	<b>11</b>
3.1	Introduction to NumPy . . . . .	11
3.2	Installing NumPy . . . . .	11
3.3	Basic Operations . . . . .	11
3.4	Array Manipulations . . . . .	12
3.5	Mathematical Functions . . . . .	12
3.6	Recap . . . . .	12
<b>4</b>	<b>Basic Operations with NumPy</b>	<b>13</b>
4.1	The @ Operator . . . . .	13
4.2	Vector-Matrix Multiplication . . . . .	13
4.3	Matrix-Matrix Multiplication . . . . .	13
4.4	Sum Across Rows . . . . .	14
4.5	Finding the Argument of the Maximum Value . . . . .	14

4.6	Recap . . . . .	14
<b>5</b>	<b>Advanced Operations with NumPy</b>	<b>15</b>
5.1	Broadcasting . . . . .	15
5.2	Multidimensional Arrays . . . . .	15
5.2.1	Creating Multidimensional Arrays . . . . .	15
5.2.2	Manipulating Multidimensional Arrays . . . . .	16
5.2.3	Accessing Elements in Multidimensional Arrays . .	16
5.3	Random Sampling . . . . .	16
5.3.1	Uniform Random Numbers . . . . .	16
5.3.2	Random Integers . . . . .	16
5.3.3	Normal Distribution . . . . .	16
5.4	Recap . . . . .	17
<b>6</b>	<b>Images as Matrices</b>	<b>19</b>
6.1	Introduction . . . . .	19
6.2	Black and White Images . . . . .	19
6.2.1	Creating a Black and White Image . . . . .	19
6.2.2	Displaying the Image . . . . .	19
6.3	RGB Images . . . . .	20
6.3.1	Creating an RGB Image . . . . .	20
6.3.2	Displaying the RGB Image . . . . .	20
6.4	Manipulating Images . . . . .	20
6.4.1	Inverting a Black and White Image . . . . .	20
6.4.2	Changing Color Channels in RGB Image . . . . .	21
6.5	Recap . . . . .	21
<b>7</b>	<b>Image Indexing</b>	<b>23</b>
7.1	Introduction . . . . .	23
7.2	Cropping Images . . . . .	23
7.2.1	Cropping a Black and White Image . . . . .	23
7.2.2	Cropping an RGB Image . . . . .	23
7.3	Selecting Color Channels . . . . .	24
7.3.1	Selecting the Red Channel . . . . .	24
7.3.2	Creating a Green Monochrome Image . . . . .	24
7.4	Combining Cropping and Channel Selection . . . . .	24

7.5	Recap . . . . .	25
<b>8</b>	<b>Loading, Resizing, and Cropping Images</b>	<b>27</b>
8.1	Introduction . . . . .	27
8.2	Loading an Image from a URL . . . . .	27
8.3	Resizing Images . . . . .	27
8.4	Cropping Images with Indexing . . . . .	28
8.5	Combining Loading, Resizing, and Cropping . . . . .	28
8.6	Recap . . . . .	29
<b>9</b>	<b>Reshaping Images into Vectors and Vice Versa</b>	<b>31</b>
9.1	Introduction . . . . .	31
9.2	Loading MNIST Data with PyTorch . . . . .	31
9.3	Visualizing an Image from MNIST . . . . .	31
9.4	Reshaping an Image into a Vector . . . . .	32
9.5	Reshaping a Vector back into an Image . . . . .	32
9.6	Applications and Insights . . . . .	32
9.7	Recap . . . . .	33
<b>10</b>	<b>Indexing Image Tensors</b>	<b>35</b>
10.1	Introduction . . . . .	35
10.2	Loading MNIST Data with PyTorch . . . . .	35
10.3	Indexing Specific Pixels . . . . .	35
10.4	Cropping Images . . . . .	35
10.5	Creating Random Images . . . . .	36
10.6	Indexing with Conditional Statements . . . . .	36
10.7	Advanced Indexing Techniques . . . . .	36
10.8	Recap . . . . .	37
<b>11</b>	<b>Working with Multi-Channel Images</b>	<b>39</b>
11.1	Introduction . . . . .	39
11.2	Creating a Multi-Channel Image . . . . .	39
11.3	Loading a Real RGB Image . . . . .	39
11.4	Channel Separation . . . . .	40
11.5	Channel Manipulation . . . . .	41
11.6	Recap . . . . .	41

- 12 Image Transformations 43**
  - 12.1 Introduction . . . . . 43
  - 12.2 Loading an Image . . . . . 43
  - 12.3 Resizing Images . . . . . 43
  - 12.4 Rotating Images . . . . . 44
  - 12.5 Flipping Images . . . . . 44
  - 12.6 Combining Transformations . . . . . 44
  - 12.7 Recap . . . . . 45
  
- III Deep Learning with PyTorch 47**
  
- 13 Introduction to PyTorch 49**
  - 13.1 What is PyTorch? . . . . . 49
  - 13.2 Installing PyTorch . . . . . 49
  - 13.3 Tensors . . . . . 49
  - 13.4 Autograd: Automatic Differentiation . . . . . 50
  - 13.5 Building a Simple Neural Network . . . . . 50
  - 13.6 Recap . . . . . 51
  
- 14 Working with MNIST Dataset: From Data Loading to Basic Operations 53**
  - 14.1 Introduction . . . . . 53
  - 14.2 Importing Libraries . . . . . 53
  - 14.3 Loading the MNIST Dataset . . . . . 53
  - 14.4 Data Preprocessing . . . . . 54
  - 14.5 Exploring and Visualizing the Data . . . . . 54
  - 14.6 Reshaping the Images into Vectors . . . . . 54
  - 14.7 Matrix Multiplication with Random Weights . . . . . 55
  
- 15 Tensor Manipulations 57**
  - 15.1 Introduction . . . . . 57
  - 15.2 Finding the Best Matrix . . . . . 57
    - 15.2.1 Initialization . . . . . 57
    - 15.2.2 Iterative Search . . . . . 57
  - 15.3 Utilizing GPU with PyTorch . . . . . 58

15.3.1	Data Conversion . . . . .	58
15.3.2	Advanced Iterative Search . . . . .	58
15.3.3	Multiple Random Matrices . . . . .	58
15.4	Testing . . . . .	59
15.5	Recap . . . . .	59
<b>16</b>	<b>Visualizing and Building the Model</b>	<b>61</b>
16.1	Introduction . . . . .	61
16.2	Data Handling and Visualization . . . . .	61
16.2.1	Converting Data to GPU . . . . .	61
16.2.2	Plotting Functions . . . . .	61
16.2.3	One-Hot Encoding . . . . .	62
16.2.4	MNIST Data Loading . . . . .	62
16.3	Building the Model . . . . .	62
16.3.1	Softmax and Cross-Entropy . . . . .	62
16.3.2	Truncated Normal Function . . . . .	63
16.3.3	Accuracy Calculation . . . . .	63
16.3.4	Batch Retrieval . . . . .	63
16.3.5	Model Definition . . . . .	64
16.3.6	Plotting and Logging . . . . .	64
16.4	Recap . . . . .	64
<b>17</b>	<b>Training the Simple Linear Model on MNIST</b>	<b>67</b>
17.1	Introduction . . . . .	67
17.2	Setting Up the Experiment . . . . .	67
17.2.1	Initializing Wandb . . . . .	67
17.2.2	Setting Hyperparameters . . . . .	67
17.3	Model Initialization and Optimization . . . . .	67
17.3.1	Initializing Weights . . . . .	67
17.3.2	Setting Up the Optimizer . . . . .	68
17.4	Training Loop . . . . .	68
17.5	Recap . . . . .	69
<b>18</b>	<b>Training a Fully Connected Neural Network on MNIST</b>	<b>71</b>
18.1	Introduction . . . . .	71
18.2	Model Definition . . . . .	71

18.2.1	Gradient Step Function . . . . .	71
18.2.2	ReLU Activation Function . . . . .	71
18.2.3	Model Architectures . . . . .	71
18.3	Training the Model . . . . .	72
18.3.1	Training with Manual Gradient Descent . . . . .	72
18.3.2	Training with the Adam Optimizer . . . . .	73
18.4	Recap . . . . .	73
<b>19</b>	<b>Building an Autoencoder for Fashion MNIST</b>	<b>75</b>
19.1	Introduction . . . . .	75
19.2	Utility Functions . . . . .	75
19.2.1	Truncated Normal Random Numbers . . . . .	75
19.2.2	GPU Functions . . . . .	75
19.3	Data Loading and Preprocessing . . . . .	76
19.4	Building an Autoencoder . . . . .	76
19.4.1	Model Definition . . . . .	76
19.4.2	Training the Autoencoder . . . . .	76
19.4.3	Testing the Autoencoder . . . . .	77
19.5	Recap . . . . .	77
<b>20</b>	<b>Convolutional Autoencoder (CAE) for Fashion MNIST</b>	<b>79</b>
20.1	Introduction . . . . .	79
20.2	Reshaping and Normalizing Data . . . . .	79
20.3	Encoder Definition . . . . .	79
20.4	Decoder Definition . . . . .	80
20.5	Training the Convolutional Autoencoder . . . . .	80
20.6	Evaluation and Visualization . . . . .	81
20.7	Recap . . . . .	81
<b>21</b>	<b>Generative Adversarial Network (GAN) on Fashion MNIST</b>	<b>83</b>
21.1	Introduction . . . . .	83
21.2	Preprocessing Data for Class Label $n$ . . . . .	83
21.3	Moving Data to GPU . . . . .	83
21.4	Discriminator Definition . . . . .	84
21.5	Generator Definition . . . . .	84
21.6	Training the GAN . . . . .	84



21.7 Visualizing Generated Images . . . . .	85
21.8 Recap . . . . .	85
<b>22 Convolutional Neural Networks</b>	<b>87</b>
22.1 Convolution . . . . .	87
22.1.1 Introduction . . . . .	87
22.1.2 Loading an Image . . . . .	87
22.1.3 Extracting the Coin Pattern . . . . .	87
22.1.4 Preprocessing the Image and Coin . . . . .	87
22.1.5 Performing Convolution . . . . .	88
22.1.6 Visualizing the Coin Detection . . . . .	88
22.1.7 Recap . . . . .	88
22.2 Applying a Convolutional Filter . . . . .	89
22.2.1 Introduction . . . . .	89
22.2.2 Loading and Preprocessing the Image . . . . .	89
22.2.3 Creating Random Filters . . . . .	89
22.2.4 Transposing the Image . . . . .	90
22.2.5 Applying the Convolutional Filter . . . . .	90
22.2.6 Recap . . . . .	90
22.3 Applying Various Convolution Filters . . . . .	91
22.3.1 Introduction . . . . .	91
22.3.2 Loading and Preprocessing the Image . . . . .	91
22.3.3 Applying Handcrafted Filters . . . . .	91
22.3.4 Applying Random Filters . . . . .	91
22.3.5 Implementing a Custom Convolution Function . . . . .	92
22.3.6 Experimenting with Various Random Filters . . . . .	92
22.3.7 Recap . . . . .	93
22.4 Comparing Performance of Different Convolution Methods . . . . .	93
22.4.1 Introduction . . . . .	93
22.4.2 Using a Custom-Defined Function . . . . .	93
22.4.3 Using SciPy's Convolution Function . . . . .	94
22.4.4 Using GPU Acceleration with PyTorch . . . . .	94
22.4.5 Recap . . . . .	95

<b>23 Image Classification Using Pre-trained AlexNet</b>	<b>97</b>
23.0.1 Introduction . . . . .	97
23.0.2 Defining the AlexNet Model . . . . .	97
23.0.3 Loading the Class Labels . . . . .	97
23.0.4 Image Preprocessing . . . . .	98
23.0.5 Loading and Classifying an Image . . . . .	98
23.0.6 Recap . . . . .	98
23.1 Exploring and Manipulating AlexNet Layers . . . . .	99
23.1.1 Introduction . . . . .	99
23.1.2 Extracting Weights from AlexNet . . . . .	99
23.1.3 Loading Weights . . . . .	99
23.1.4 Applying Layers Manually . . . . .	100
23.1.5 Recap . . . . .	101
23.2 Understanding ReLU and Max Pooling . . . . .	102
23.2.1 ReLU: Rectified Linear Unit . . . . .	102
23.2.2 Max Pooling . . . . .	102
23.2.3 Recap . . . . .	103
23.3 Training AlexNet with Custom Dataset . . . . .	104
23.3.1 Mounting Google Drive . . . . .	104
23.3.2 Loading and Pre-processing Data . . . . .	104
23.3.3 AlexNet Architecture . . . . .	105
23.3.4 Training Process . . . . .	105
23.3.5 Analysis and Summary . . . . .	107
23.4 Model Saving, Loading, and Inference . . . . .	107
23.4.1 Saving the Trained Model . . . . .	107
23.4.2 Loading the Trained Model . . . . .	107
23.4.3 Preprocessing a New Image . . . . .	108
23.4.4 Performing Inference . . . . .	108
23.4.5 Analysis and Summary . . . . .	109
<b>24 Recap</b>	<b>111</b>
24.1 Reflection on Key Concepts . . . . .	111
24.2 Future Perspectives . . . . .	111
24.3 Final Thoughts . . . . .	112

# Chapter 1

## Introduction to Data Science with PyTorch

### 1.1 The Power of Data Science

Data Science is an interdisciplinary field that utilizes scientific methods, processes, algorithms, and systems to extract knowledge and insights from structured and unstructured data. It's a field that has rapidly become essential in various industries, including finance, healthcare, marketing, and technology.

### 1.2 Tools for the Modern Data Scientist

In this book, we will focus on three powerful tools that every data scientist should know:

- **NumPy:** A library for the Python programming language, adding support for large, multi-dimensional arrays and matrices, along with mathematical functions to operate on these arrays.
- **PyTorch:** An open-source machine learning library used for applications such as computer vision and natural language processing. It is known for being both flexible and efficient.
- **Matplotlib:** A plotting library for creating static, interactive, and animated visualizations in Python. It is widely used for making complex plots from data.

### 1.3 A Glimpse of Python Code

Here's a simple example of Python code using NumPy to demonstrate how you can perform array operations:

```
import numpy as np
```

```
# Create a 1D array
arr = np.array([1, 2, 3, 4, 5])

# Perform element-wise addition
result = arr + 5

print(result)  # Output will be [6, 7, 8, 9, 10]
```

## 1.4 Structure of This Book

This book is structured to take you on a journey through data science, starting from the basics and working towards more advanced concepts and applications. We'll explore NumPy for numerical computations, PyTorch for deep learning, and Matplotlib for data visualization.

1. **Part I:** Fundamentals of Data Science
2. **Part II:** NumPy for Numerical Analysis
3. **Part III:** Deep Learning with PyTorch
4. **Part IV:** Visualizing Data with Matplotlib
5. **Part V:** Advanced Techniques and Applications

## 1.5 Recap

By the end of this book, you will have a solid understanding of how to leverage NumPy, PyTorch, and Matplotlib in your data science projects. Whether you are a student, professional, or hobbyist, the knowledge and skills acquired here will undoubtedly prove valuable in your data science journey. Let's get started!

## Part I

# Fundamentals of Data Science



## Chapter 2

### Understanding Data

#### 2.1 What is Data?

Data refers to facts, statistics, or information that are represented in various forms. It can be numbers, words, measurements, observations, or even pictures. In the realm of data science, data is used to create models, make predictions, and guide decisions.

#### 2.2 Types of Data

In data science, we generally categorize data into the following types:

- **Numerical Data:** Quantitative data that can be measured, such as height, weight, temperature, etc.
- **Categorical Data:** Qualitative data that can be categorized into groups or classes, such as gender, blood type, or education level.
- **Ordinal Data:** Similar to categorical data but with a meaningful order, such as rankings or ratings.
- **Time-Series Data:** Data collected or recorded at specific time intervals.

#### 2.3 Data Cleaning

Data cleaning is the process of detecting and correcting errors and inconsistencies in data to improve its quality. This involves handling missing values, removing duplicates, and correcting errors.

```
import pandas as pd

# Reading a dataset
```

```
data = pd.read_csv('data.csv')

# Removing duplicates
data = data.drop_duplicates()

# Filling missing values with the mean
data = data.fillna(data.mean())
```

## 2.4 Data Visualization

Data visualization is the graphical representation of information and data. With the use of visual elements like charts, graphs, and maps, data visualization tools provide an accessible way to see and understand trends, outliers, and patterns in data.

Here's an example of a simple plot using Matplotlib:

```
import matplotlib.pyplot as plt

# Sample data
x = [1, 2, 3, 4, 5]
y = [10, 20, 30, 40, 50]

# Creating the plot
plt.plot(x, y)

# Adding title and labels
plt.title('A Simple Plot')
plt.xlabel('X-axis')
plt.ylabel('Y-axis')

# Displaying the plot
plt.show()
```

## 2.5 Recap

Understanding the nature of data, its types, and how to process and visualize it is foundational to data science. In the next chapters, we will



delve into specific methods and techniques to analyze and interpret data, building the skills necessary to become proficient in this exciting field.



# Part II

## NumPy for Numerical Analysis



## Chapter 3

# Getting Started with NumPy

### 3.1 Introduction to NumPy

NumPy, short for Numerical Python, is a library for the Python programming language that provides support for large, multi-dimensional arrays and matrices. It also offers a wide variety of mathematical functions to operate on these arrays, making it essential for scientific computing.

### 3.2 Installing NumPy

You can install NumPy using pip, the Python package installer. Open a terminal or command prompt and run the following command:

```
pip install numpy
```

### 3.3 Basic Operations

Once NumPy is installed, you can import it and start working with arrays. Here's an example of basic array creation and arithmetic operations:

```
import numpy as np

# Creating arrays
a = np.array([1, 2, 3])
b = np.array([4, 5, 6])

# Arithmetic operations
sum_array = a + b
product_array = a * b
```

## 3.4 Array Manipulations

NumPy provides various functions to reshape, split, and combine arrays.

```
# Reshaping an array
reshaped_array = np.reshape(a, (3, 1))

# Splitting an array
split_arrays = np.split(b, 3)

# Stacking arrays vertically
stacked_array = np.vstack((a, b))
```

## 3.5 Mathematical Functions

NumPy provides a wide range of mathematical functions that can be applied to arrays, including trigonometric, logarithmic, and statistical functions.

```
# Trigonometric function
sin_values = np.sin(a)

# Logarithmic function
log_values = np.log(b)

# Statistical function
mean_value = np.mean(a)
```

## 3.6 Recap

This chapter provides an introduction to NumPy and covers essential concepts and functions for working with arrays. As we progress through this part of the book, we'll explore more advanced features of NumPy.

Understanding NumPy is vital for anyone working in data science, as it lays the foundation for other libraries like SciPy, Pandas, and even deep learning frameworks like PyTorch.

## Chapter 4

### Basic Operations with NumPy

#### 4.1 The @ Operator

In NumPy, the "@" operator is used as an infix operator for matrix multiplication. It's a more readable and convenient way to multiply matrices compared to the 'np.dot' function.

```
import numpy as np

a = np.array([1, 2])
b = np.array([[4, 5], [6, 7]])

# Using the @ operator for matrix multiplication
result = a @ b
```

#### 4.2 Vector-Matrix Multiplication

Vector-matrix multiplication is a key operation in linear algebra. Here's how you can perform this multiplication using NumPy:

```
vector = np.array([1, 2, 3])
matrix = np.array([[4, 5], [6, 7], [8, 9]])

# Multiplying the vector and the matrix
result = vector @ matrix
```

#### 4.3 Matrix-Matrix Multiplication

Matrix-matrix multiplication can be done similarly using the "@" operator:

```
matrix1 = np.array([[1, 2], [3, 4]])
matrix2 = np.array([[5, 6], [7, 8]])
```

```
# Multiplying the matrices
result = matrix1 @ matrix2
```

## 4.4 Sum Across Rows

You can use the ‘np.sum’ function with the ‘axis’ parameter to sum across rows:

```
matrix = np.array([[1, 2, 3], [4, 5, 6]])

# Summing across rows
row_sums = np.sum(matrix, axis=1)
```

## 4.5 Finding the Argument of the Maximum Value

The ‘np.argmax’ function returns the indices of the maximum values along an axis. Here’s how you can use it:

```
array = np.array([10, 20, 30, 20])

# Finding the index of the maximum value
max_index = np.argmax(array)
```

## 4.6 Recap

This chapter has covered essential operations in NumPy, including matrix multiplications using the “@” operator, summing across rows, and finding the argument of the maximum value. These basic operations form the foundation of many algorithms and techniques in data science and machine learning.

In the following chapters, we will continue to explore the capabilities of NumPy by diving into more complex mathematical functions and methods for data manipulation.



## Chapter 5

### Advanced Operations with NumPy

#### 5.1 Broadcasting

Broadcasting is a powerful mechanism that allows NumPy to work with arrays of different shapes when performing arithmetic operations. It extends smaller arrays to match the shape of the larger array.

```
import numpy as np

a = np.array([1, 2, 3])
b = 2

# Broadcasting b across a
result = a * b

# Broadcasting can also work with higher dimensions
matrix = np.array([[1, 2], [3, 4]])
vector = np.array([10, 20])

# Broadcasting vector across matrix
result_matrix = matrix * vector
```

#### 5.2 Multidimensional Arrays

NumPy allows you to work with arrays that have more than two dimensions, enabling complex data structures.

##### 5.2.1 Creating Multidimensional Arrays

```
# Creating a 3-dimensional array
array_3d = np.array([[[1, 2], [3, 4]], [[5, 6], [7, 8]]])
```

## 5.2.2 Manipulating Multidimensional Arrays

```
# Transposing a matrix
transposed = np.transpose(matrix)

# Reshaping a 3-dimensional array
reshaped = array_3d.reshape((4, 2))
```

## 5.2.3 Accessing Elements in Multidimensional Arrays

```
# Accessing an element in a 3-dimensional array
element = array_3d[1, 0, 1]
```

## 5.3 Random Sampling

Random sampling is a core concept in statistics and machine learning. NumPy offers an extensive set of functions to generate random samples.

### 5.3.1 Uniform Random Numbers

```
# Generating random numbers uniformly between 0 and 1
uniform_randoms = np.random.rand(10)
```

### 5.3.2 Random Integers

```
# Generating random integers between 1 and 10
random_integers = np.random.randint(low=1, high=10,
    size=5)
```

### 5.3.3 Normal Distribution

```
# Generating random numbers from the normal
distribution
random_numbers = np.random.normal(loc=0, scale=1, size
    =100)
```

## 5.4 Recap

This chapter dove into advanced concepts and techniques in NumPy, including the essential idea of broadcasting, working with multidimensional arrays, and generating random samples. These techniques are pivotal in data manipulation and analysis, and they set the foundation for more complex data science and machine learning tasks.

The knowledge gained here will be instrumental as we explore other libraries and delve into specialized areas of data science in the subsequent parts of the book.



# Chapter 6

## Images as Matrices

### 6.1 Introduction

Images can be thought of as numerical matrices, where each element represents a pixel's intensity. This chapter explores the representation of images as matrices and demonstrates how to manipulate and visualize them using NumPy and Matplotlib.

### 6.2 Black and White Images

A black and white image can be represented as a 2D matrix, where each value corresponds to a pixel's intensity, ranging from 0 (black) to 255 (white).

#### 6.2.1 Creating a Black and White Image

```
import numpy as np

# Creating a simple black and white image
bw_image = np.array([[0, 255, 0], [255, 255, 255], [0,
                255, 0]])
```

#### 6.2.2 Displaying the Image

```
import matplotlib.pyplot as plt

plt.imshow(bw_image, cmap='gray')
plt.axis('off') # to turn off axes
plt.show()
```

## 6.3 RGB Images

An RGB image consists of three color channels: Red, Green, and Blue. It can be represented as a 3D matrix, where the third dimension has size 3, corresponding to the three color channels.

### 6.3.1 Creating an RGB Image

```
# Creating a simple RGB image
rgb_image = np.array([[[255, 0, 0], [0, 255, 0], [0,
    0, 255]],
                      [[0, 255, 255], [255, 0, 255],
                       [255, 255, 0]],
                      [[128, 128, 128], [255, 255,
                       255], [0, 0, 0]]])
```

### 6.3.2 Displaying the RGB Image

```
plt.imshow(rgb_image)
plt.axis('off') # to turn off axes
plt.show()
```

## 6.4 Manipulating Images

Images, being represented as matrices, can be manipulated using standard mathematical operations.

### 6.4.1 Inverting a Black and White Image

```
# Inverting a black and white image
inverted_bw_image = 255 - bw_image

plt.imshow(inverted_bw_image, cmap='gray')
plt.axis('off')
plt.show()
```

### 6.4.2 Changing Color Channels in RGB Image

```
# Swapping the Red and Green channels
swapped_channels_image = rgb_image.copy()
swapped_channels_image[:, :, 0],
    swapped_channels_image[:, :, 1] = rgb_image[:, :,
    1], rgb_image[:, :, 0]

plt.imshow(swapped_channels_image)
plt.axis('off')
plt.show()
```

## 6.5 Recap

This chapter unveiled the fascinating connection between images and matrices. By understanding images as numerical arrays, we have the power to manipulate them at the pixel level. This concept is crucial in many areas, including computer vision, image processing, and deep learning with convolutional neural networks.

In subsequent chapters, we will build on these foundations to explore more advanced image processing techniques and delve into the world of machine learning with PyTorch.





# Chapter 7

## Image Indexing

### 7.1 Introduction

Indexing is a powerful technique to access specific parts of an image or perform particular operations on specific channels. In this chapter, we will explore how to use indexing to crop images and manipulate individual color channels.

### 7.2 Cropping Images

Cropping is the act of selecting a specific region of an image. This can be achieved by selecting a subset of the image's matrix.

#### 7.2.1 Cropping a Black and White Image

```
# Selecting the middle pixel
cropped_bw_image = bw_image[1:2, 1:2]

plt.imshow(cropped_bw_image, cmap='gray')
plt.axis('off')
plt.show()
```

#### 7.2.2 Cropping an RGB Image

```
# Selecting the middle square region
cropped_rgb_image = rgb_image[0:2, 0:2, :]

plt.imshow(cropped_rgb_image)
plt.axis('off')
plt.show()
```

## 7.3 Selecting Color Channels

Working with color images often requires manipulating individual channels. Indexing allows us to isolate and operate on these channels.

### 7.3.1 Selecting the Red Channel

```
# Selecting only the Red channel
red_channel = rgb_image[:, :, 0]

plt.imshow(red_channel, cmap='gray')
plt.axis('off')
plt.show()
```

### 7.3.2 Creating a Green Monochrome Image

```
# Zeroing out the Red and Blue channels
green_monochrome_image = rgb_image.copy()
green_monochrome_image[:, :, 0] = 0
green_monochrome_image[:, :, 2] = 0

plt.imshow(green_monochrome_image)
plt.axis('off')
plt.show()
```

## 7.4 Combining Cropping and Channel Selection

```
# Cropping the image and selecting the Blue channel
cropped_blue_channel = rgb_image[0:2, 1:3, 2]

plt.imshow(cropped_blue_channel, cmap='gray')
plt.axis('off')
plt.show()
```

## 7.5 Recap

Understanding images as matrices unlocks a world of possibilities for manipulation and analysis. This chapter demonstrated the ease with which we can crop images and select individual color channels using indexing.

The principles and techniques explored here form the building blocks of more complex image processing and computer vision tasks, such as segmentation, feature extraction, and object detection. They also set the stage for understanding convolutional neural networks, a critical component in modern deep learning.



## Chapter 8

### Loading, Resizing, and Cropping Images

#### 8.1 Introduction

Real-world applications often require loading images from external sources, resizing them to specific dimensions, and cropping parts of interest. This chapter will guide you through these essential image manipulation techniques using Python libraries, including ‘skimage’ and ‘matplotlib’.

#### 8.2 Loading an Image from a URL

With ‘skimage’, loading an image from a URL is straightforward. The ‘imread’ function can read images from local paths and URLs.

```
from skimage.io import imread
import matplotlib.pyplot as plt

url = 'https://tinyurl.com/RGBImageCat'
im = imread(url)

plt.figure(figsize=(20,10))
plt.imshow(im)
plt.axis('off')
plt.show()
```

#### 8.3 Resizing Images

Resizing images is often necessary to standardize dimensions. We can use the ‘skimage.transform.resize’ method.

```
from skimage.transform import resize

# Resizing the image to 100x100 pixels
resized_image = resize(im, (100, 100))
```

```
plt.imshow(resized_image)
plt.axis('off')
plt.show()
```

## 8.4 Cropping Images with Indexing

As previously discussed, images can be cropped using indexing. With a loaded image, we can select a region of interest.

```
# Cropping the central part of the image
height, width, _ = im.shape
cropped_image = im[height//4:3*height//4, width//4:3*
    width//4]

plt.imshow(cropped_image)
plt.axis('off')
plt.show()
```

## 8.5 Combining Loading, Resizing, and Cropping

We can combine these techniques to process an image effectively. Let's load an image from a URL, resize it, and then crop a specific region.

```
# Load the image
image = imread(url)

# Resize it
resized_image = resize(image, (200, 200))

# Crop the central part
cropped_image = resized_image[50:150, 50:150]

plt.imshow(cropped_image)
plt.axis('off')
plt.show()
```

## 8.6 Recap

This chapter provides a hands-on guide to loading, resizing, and cropping images using ‘skimage’, equipping you with fundamental skills for image manipulation. These techniques are foundational in various applications, such as data preparation for machine learning models, image analysis, and computer vision tasks.

As we move forward, we will build on these skills to explore more complex image processing and machine learning concepts, unlocking the potential to create advanced computer vision applications.





## Chapter 9

# Reshaping Images into Vectors and Vice Versa

## 9.1 Introduction

Understanding the dimensions and shapes of images and vectors is essential when working with image data. In this chapter, we'll explore how to reshape images into vectors and vice versa, using the MNIST dataset loaded with PyTorch. We will also print and discuss the shapes at each step to provide insights into the transformations.

## 9.2 Loading MNIST Data with PyTorch

```
from torchvision.datasets import MNIST
from torchvision.transforms import ToTensor

train_data = MNIST(root='data', train=True, download=
    True, transform=ToTensor())
test_data = MNIST(root='data', train=False, transform=
    ToTensor())

train_image, train_label = train_data[0]
print(train_image.shape)  # Output: torch.Size([1, 28,
    28])
```

The MNIST images have dimensions of  $1 \times 28 \times 28$ , where the first dimension represents the number of color channels (1 for grayscale), and the next two dimensions represent the height and width of the image.

## 9.3 Visualizing an Image from MNIST

```
plt.imshow(train_image.squeeze().numpy(), cmap='gray')
plt.axis('off')
plt.show()
```

## 9.4 Reshaping an Image into a Vector

```
image_vector = train_image.reshape(-1)
print(image_vector.shape)  # Output: torch.Size([784])

plt.plot(image_vector)
plt.title('Image Vector')
plt.show()
```

The image has been reshaped into a vector of shape 784, by multiplying the height and width dimensions  $28 \times 28 = 784$ .

## 9.5 Reshaping a Vector back into an Image

```
reshaped_image = image_vector.reshape(train_image.
    shape[1:])
print(reshaped_image.shape)  # Output: torch.Size([28,
    28])

plt.imshow(reshaped_image.numpy(), cmap='gray')
plt.axis('off')
plt.show()
```

Here, the vector has been reshaped back into its original  $28 \times 28$  dimensions, representing the height and width of the image.

## 9.6 Applications and Insights

Understanding and manipulating the dimensions of images and vectors are critical skills in machine learning and computer vision. The reshaping process demonstrated here is often employed in various applications, including feeding image data into neural networks.

## 9.7 Recap

This chapter has elucidated the reshaping process with detailed print-outs of shapes and an explanation of dimensions, using PyTorch and the MNIST dataset. These fundamental operations form the basis for more advanced techniques in handling images, to be explored in subsequent chapters.



## Chapter 10

### Indexing Image Tensors

#### 10.1 Introduction

Indexing tensors allows us to access specific elements or sub-tensors within a tensor. This is crucial when working with image data, whether for cropping, segmentation, or modifying specific parts of an image. In this chapter, we will explore different techniques for indexing image tensors using the MNIST dataset and randomly generated values.

#### 10.2 Loading MNIST Data with PyTorch

```
from torchvision.datasets import MNIST
from torchvision.transforms import ToTensor

train_data = MNIST(root='data', train=True, download=
    True, transform=ToTensor())
train_image, train_label = train_data[0]
```

#### 10.3 Indexing Specific Pixels

We can access a specific pixel by providing its coordinates.

```
pixel_value = train_image[0, 14, 14]
print(pixel_value) # Prints the value of the pixel at
    the center of the image
```

#### 10.4 Cropping Images

Cropping can be performed by indexing a specific range of pixels.

```
cropped_image = train_image[:, 10:20, 10:20]

plt.imshow(cropped_image.squeeze().numpy(), cmap='gray')
plt.axis('off')
plt.show()
```

## 10.5 Creating Random Images

We can create a random image tensor and apply the same indexing techniques.

```
import torch

random_image = torch.rand(1, 28, 28)
plt.imshow(random_image.squeeze().numpy(), cmap='gray')
plt.axis('off')
plt.show()
```

## 10.6 Indexing with Conditional Statements

We can apply conditions to create masks for indexing.

```
mask = random_image > 0.5
masked_image = random_image * mask

plt.imshow(masked_image.squeeze().numpy(), cmap='gray')
plt.axis('off')
plt.show()
```

## 10.7 Advanced Indexing Techniques

Advanced indexing techniques can be used to manipulate and transform images in complex ways, such as flipping, rotating, or custom segmenta-

tions.

## 10.8 Recap

Indexing is a powerful tool for manipulating and analyzing image tensors. Through various examples with the MNIST dataset and random values, this chapter has demonstrated how to access and modify specific parts of images. These techniques are foundational in image processing and computer vision, enabling more complex operations and analysis.





# Chapter 11

## Working with Multi-Channel Images

### 11.1 Introduction

Images are often represented in multiple channels, such as Red, Green, and Blue (RGB) channels in a color image. Understanding how to manipulate and analyze multi-channel images is essential in many fields, including computer vision, machine learning, and digital media. In this chapter, we will explore how to work with these channels using Python.

### 11.2 Creating a Multi-Channel Image

```
import numpy as np

# Creating a synthetic RGB image
red_channel = 255 * np.ones((100, 100), dtype=np.uint8)
green_channel = np.zeros((100, 100), dtype=np.uint8)
blue_channel = np.zeros((100, 100), dtype=np.uint8)

image_rgb = np.stack([red_channel, green_channel,
                      blue_channel], axis=2)

plt.imshow(image_rgb)
plt.axis('off')
plt.show()
```

Here, we create a synthetic red image by stacking three 100x100 channels. The red channel is filled with the maximum value (255), while the green and blue channels are filled with zeros.

### 11.3 Loading a Real RGB Image

```
from skimage.io import imread

url = 'https://tinyurl.com/RGBImageCat'
image_real = imread(url)

plt.imshow(image_real)
plt.axis('off')
plt.show()
```

## 11.4 Channel Separation

We can separate the channels of an RGB image for individual analysis.

```
red, green, blue = image_real[:, :, 0], image_real[:, :, 1], image_real[:, :, 2]

plt.figure(figsize=(15,5))
plt.subplot(1,3,1)
plt.imshow(red, cmap='gray')
plt.title('Red Channel')
plt.axis('off')

plt.subplot(1,3,2)
plt.imshow(green, cmap='gray')
plt.title('Green Channel')
plt.axis('off')

plt.subplot(1,3,3)
plt.imshow(blue, cmap='gray')
plt.title('Blue Channel')
plt.axis('off')

plt.show()
```

## 11.5 Channel Manipulation

Channels can be manipulated to achieve various effects, such as color balancing, filtering, or creating artistic effects.

```
# Increasing the intensity of the green channel
image_modified = image_real.copy()
image_modified[:, :, 1] = image_modified[:, :, 1] * 1.5 #
    Increase green

plt.imshow(image_modified)
plt.axis('off')
plt.show()
```

## 11.6 Recap

Multi-channel images open up a wide range of possibilities for analysis, manipulation, and interpretation. This chapter has introduced fundamental techniques for handling RGB images, which can be extended to other color spaces and more complex operations. Understanding these principles is vital for anyone working with image data, from data scientists to graphic designers.



# Chapter 12

## Image Transformations

### 12.1 Introduction

Image transformations play an essential role in pre-processing, data augmentation, and various applications in computer vision and graphics. This chapter will focus on resizing, rotating, and flipping images, fundamental techniques widely used in practice.

### 12.2 Loading an Image

Let's start by loading an image using the same method as in previous chapters.

```
from skimage.io import imread

url = 'https://example.com/image.jpg'
image = imread(url)

plt.imshow(image)
plt.axis('off')
plt.show()
```

### 12.3 Resizing Images

Resizing is often necessary to normalize images to a consistent dimension or to reduce computational requirements.

```
from skimage.transform import resize

resized_image = resize(image, (50, 50))

plt.imshow(resized_image)
plt.axis('off')
```

```
plt.show()
```

## 12.4 Rotating Images

Rotation can be used for data augmentation or to correct the orientation of an image.

```
from skimage.transform import rotate

rotated_image = rotate(image, angle=45)

plt.imshow(rotated_image)
plt.axis('off')
plt.show()
```

## 12.5 Flipping Images

Flipping can create reflections of the image, often used for data augmentation.

```
from skimage.transform import hflip

flipped_image = hflip(image)

plt.imshow(flipped_image)
plt.axis('off')
plt.show()
```

## 12.6 Combining Transformations

Transformations can be combined to create complex effects or extensive data augmentations.

```
# Combining resize, rotate, and flip
combined_image = hflip(rotate(resize(image, (100, 100)
), angle=30))
```

```
plt.imshow(combined_image)
plt.axis('off')
plt.show()
```

## 12.7 Recap

Image transformations are fundamental in image processing and provide powerful tools for manipulation and augmentation. By understanding how to resize, rotate, and flip images, you can prepare data for machine learning models, create artistic effects, or correct and enhance digital media. These concepts serve as building blocks for more advanced techniques and applications in the field of computer vision.





## Part III

# Deep Learning with PyTorch



## Chapter 13

# Introduction to PyTorch

### 13.1 What is PyTorch?

PyTorch is an open-source machine learning library developed by Facebook's AI Research lab. It is known for its flexibility, efficiency, and being highly similar to Python in its design. PyTorch has become one of the preferred frameworks for deep learning research and development.

### 13.2 Installing PyTorch

You can install PyTorch by following the instructions on the official PyTorch website. The installation command varies based on your system and whether you have a compatible GPU. Here's an example command for a standard CPU installation:

```
pip install torch torchvision
```

### 13.3 Tensors

Tensors are multi-dimensional arrays and are the fundamental building blocks in PyTorch. They can be created and manipulated much like NumPy arrays:

```
import torch

# Creating a tensor
a = torch.tensor([1, 2, 3])

# Basic operations
b = a + 5
```

## 13.4 Autograd: Automatic Differentiation

PyTorch provides automatic differentiation, which is essential for training neural networks. The ‘autograd’ package computes derivatives and gradients automatically:

```
x = torch.tensor(2.0, requires_grad=True)
y = x**2
y.backward() # Computes the gradient
gradient = x.grad # gradient will be 4
```

## 13.5 Building a Simple Neural Network

Here’s an example of a basic feed-forward network with one hidden layer:

```
import torch
import torch.nn as nn
import torch.nn.functional as F

input_size = 3
hidden_size = 5
output_size = 1

# Define the weights and biases
W1 = torch.randn(input_size, hidden_size,
                  requires_grad=True)
b1 = torch.randn(hidden_size, requires_grad=True)
W2 = torch.randn(hidden_size, output_size,
                  requires_grad=True)
b2 = torch.randn(output_size, requires_grad=True)

# Example input
x = torch.tensor([1.0, 2.0, 3.0])

# Forward pass
hidden = F.relu(x @ W1 + b1)
output = hidden @ W2 + b2
```

This code creates a simple neural network using basic PyTorch operations and functions. The ‘@’ symbol represents matrix multiplication, and ‘F.relu’ is the ReLU activation function from the ‘torch.nn.functional’ module.

## 13.6 Recap

This chapter has introduced the core concepts of PyTorch, including tensors, automatic differentiation, and neural network construction. In the following chapters, we’ll delve into more advanced topics like convolutional neural networks, recurrent neural networks, and techniques for training and fine-tuning models.

PyTorch offers a rich ecosystem of tools and libraries that cater to different research and production needs. Understanding the fundamentals of PyTorch will open doors to an array of applications in machine learning and artificial intelligence.



## Chapter 14

# Working with MNIST Dataset: From Data Loading to Basic Operations

### 14.1 Introduction

In this chapter, we will explore the code that works with the MNIST dataset, one of the most well-known datasets in the machine learning community. This dataset consists of 60,000 training and 10,000 testing images, each a grayscale image of a handwritten digit (from 0 to 9). We'll cover loading the data, visualizing the images, flattening and reshaping them, and performing some basic matrix operations.

### 14.2 Importing Libraries

```
import torch
import numpy as np
import torchvision
import torchvision.datasets as datasets
import matplotlib.pyplot as plt
```

Here we import the essential libraries: PyTorch, NumPy for numerical computations, torchvision for loading datasets, and Matplotlib for plotting.

### 14.3 Loading the MNIST Dataset

```
train_set = datasets.MNIST('./data', train=True,
                           download=True)
test_set = datasets.MNIST('./data', train=False,
                          download=True)
```

These lines load the MNIST dataset for training and testing. The images will be downloaded if they are not present in the specified directory.

## 14.4 Data Preprocessing

```
X = train_set.data.numpy()
X_test = test_set.data.numpy()
Y = train_set.targets.numpy()
Y_test = test_set.targets.numpy()
X = X[:,None,:,:]/255
X_test = X_test[:,None,:,:]/255
```

Here, we convert the data and targets to NumPy arrays and divide the pixel values by 255 to normalize them.

## 14.5 Exploring and Visualizing the Data

```
plt.imshow(X[0,0,:,:], cmap='gray')
for i in range(10):
    plt.imshow(X[i,0,:,:], cmap='gray')
    plt.title(str(Y[i]))
plt.show()
```

This code visualizes the first 10 images in the dataset and prints the corresponding labels as titles.

## 14.6 Reshaping the Images into Vectors

```
x = X[0,0,:,:].flatten()
X = np.reshape(X, (X.shape[0], X.shape[2]*X.shape[3]))
X_test = np.reshape(X_test, (X_test.shape[0], X_test.
    shape[2]*X_test.shape[3]))
```

Here, we flatten the images from 28x28 matrices to 784-length vectors to make them suitable for matrix multiplication.



## 14.7 Matrix Multiplication with Random Weights

```
m = np.random.standard_normal((10,784))
y = m@X
y = np.argmax(y, axis=0)
acc = np.sum(y == Y)/len(Y)
print(acc)
```

This code snippet randomly initializes a matrix of weights and performs matrix multiplication on the dataset. Then, it calculates the argmax and computes a basic accuracy measure.



# Chapter 15

## Tensor Manipulations

### 15.1 Introduction

In this chapter, we delve into the code that iteratively seeks an optimal alignment between a randomly generated matrix and the MNIST dataset. We'll also learn about the usage of GPUs for computations with PyTorch. These concepts represent fundamental techniques that provide deeper insights into data alignment and matrix operations.

### 15.2 Finding the Best Matrix

#### 15.2.1 Initialization

```
m_best = 0
acc_best = 0
```

Here, we initialize the variables that will hold the best matrix ‘m best’ and its corresponding accuracy ‘acc best’.

#### 15.2.2 Iterative Search

```
for i in range(1000):
    m = 0.1*np.random.standard_normal((10,784))
    y = m@X
    y = np.argmax(y, axis=0)
    acc = np.sum(y == y_ans)/len(Y)
    if acc > acc_best:
        print(acc)
        m_best = m
        acc_best = acc
```

This loop iterates 1000 times, each time creating a new random matrix ‘m’ and calculating its accuracy ‘acc’ by performing matrix multiplication

with the data ‘X’ and comparing the result with the actual labels ‘Y’. If the current accuracy is higher than the best accuracy so far, it updates ‘m best’ and ‘acc best’.

## 15.3 Utilizing GPU with PyTorch

### 15.3.1 Data Conversion

```
def GPU(data):  
    return torch.tensor(data, requires_grad=True,  
                        dtype=torch.float, device=torch.device('cuda'))  
  
def GPU_data(data):  
    return torch.tensor(data, requires_grad=False,  
                        dtype=torch.float, device=torch.device('cuda'))  
X = GPU_data(X)  
Y = GPU_data(Y)  
X_test = GPU_data(X_test)  
Y_test = GPU_data(Y_test)
```

Here, we define two functions, ‘GPU’ and ‘GPU data’, to convert the data into PyTorch tensors and transfer them to the GPU. This step allows for faster computations.

### 15.3.2 Advanced Iterative Search

The next section of code is an extended version of the iterative search with additional complexity. It involves multiple iterations and manipulations with random matrices and steps, applying them on the GPU.

### 15.3.3 Multiple Random Matrices

```
N = 100  
M = GPU_data(np.random.rand(N,10,784))  
m_best = 0  
acc_best = 0  
step = 0.000000000001
```

```
for i in range(1000000):
    y = torch.argmax((M@X), axis=1)
    score = ((y == Y).sum(1)/len(Y))
    s = torch.argsort(score,descending=True)
    M = M[s]
    M[1:] = M[0]
    M[N//2:] = 0
    M[1:] += step*GPU_data(np.random.rand(N-1,10,784))
    acc = score[s][0].item()
    if acc > acc_best:
        m_best = M[0]
        acc_best = acc
    print(i,acc)
```

This code snippet introduces a more complex search algorithm, involving multiple random matrices, sorting, and incremental updates. The aim is the same: to find the matrix that best aligns with the dataset.

## 15.4 Testing

```
y_test = torch.argmax((M@X_test), axis=1)
score = ((y_test == Y_test).sum(1)/len(Y_test))
s = torch.argsort(score,descending=True)
acc_test = score[s][0].item()
acc_test
```

Finally, this section tests the previously found matrix on the test dataset and prints the accuracy.

## 15.5 Recap

This chapter has explained the code that demonstrates the search for an optimal matrix through iterative methods and introduced GPU computations with PyTorch. Although this approach is rather unconventional in the context of machine learning, the underlying principles are highly educational and give a glimpse into the possibilities of matrix operations, alignment, and optimization.



## Chapter 16

# Visualizing and Building the Model

## 16.1 Introduction

In this chapter, we'll explore the code that provides all the necessary building blocks to run the demo. We will delve into functions for handling the data, visualizing the data, modeling the system, and leveraging the wandb library for logging the performance.

## 16.2 Data Handling and Visualization

### 16.2.1 Converting Data to GPU

```
def GPU(data):  
    return torch.tensor(data, requires_grad=True,  
                        dtype=torch.float, device=torch.device('cuda'))  
  
def GPU_data(data):  
    return torch.tensor(data, requires_grad=False,  
                        dtype=torch.float, device=torch.device('cuda'))
```

These functions enable conversion of data into PyTorch tensors and transfer to the GPU for efficient computation.

### 16.2.2 Plotting Functions

```
def plot(x):  
    if type(x) == torch.Tensor :  
        x = x.cpu().detach().numpy()  
    fig, ax = plt.subplots()  
    im = ax.imshow(x, cmap = 'gray')  
    ax.axis('off')  
    fig.set_size_inches(10, 10)  
    plt.show()
```

```
def montage_plot(x):  
    x = np.pad(x, pad_width=((0, 0), (1, 1), (1, 1)),  
               mode='constant', constant_values=0)  
    plot(montage(x))
```

These functions are used for plotting various forms of data, including grayscale images and cool colored images. The ‘montage plot’ function includes padding and utilizes the montage function from the skimage library.

### 16.2.3 One-Hot Encoding

```
def one_hot(y):  
    y2 = GPU_data(torch.zeros((y.shape[0], 10)))  
    for i in range(y.shape[0]):  
        y2[i, int(y[i])] = 1  
    return y2
```

This function creates one-hot encoded vectors for categorical labels, a common practice in machine learning.

### 16.2.4 MNIST Data Loading

```
train_set = datasets.MNIST('./data', train=True,  
                           download=True)  
test_set = datasets.MNIST('./data', train=False,  
                           download=True)
```

This code snippet loads the MNIST dataset, a large database of handwritten digits that is widely used for training various image processing systems.

## 16.3 Building the Model

### 16.3.1 Softmax and Cross-Entropy



```
def softmax(x):  
    s1 = torch.exp(x - torch.max(x,1)[0][:,None])  
    s = s1 / s1.sum(1)[: ,None]  
    return s  
def cross_entropy(outputs, labels):  
    return -torch.sum(softmax(outputs).log()[range(  
        outputs.size()[0]), labels.long()])/outputs.  
        size()[0]
```

Here, the softmax function is implemented, which converts a vector into a probability distribution. The cross-entropy function then calculates the loss between the predicted probabilities and the actual labels.

### 16.3.2 Truncated Normal Function

```
def Truncated_Normal(size):  
    u1 = torch.rand(size)*(1-np.exp(-2)) + np.exp(-2)  
    u2 = torch.rand(size)  
    z = torch.sqrt(-2*torch.log(u1)) * torch.cos(2*np.  
        .pi*u2)  
    return z
```

This function generates values from a truncated normal distribution, a normal distribution that is limited to a range.

### 16.3.3 Accuracy Calculation

```
def acc(out,y):  
    with torch.no_grad():  
        return (torch.sum(torch.max(out,1)[1] == y).  
            item())/y.shape[0]
```

This function calculates the accuracy of predictions against the true labels.

### 16.3.4 Batch Retrieval

```
def get_batch(mode):
    b = c.b
    if mode == "train":
        r = np.random.randint(X.shape[0] - b)
        x = X[r:r+b,:]
        y = Y[r:r+b]
    elif mode == "test":
        r = np.random.randint(X_test.shape[0] - b)
        x = X_test[r:r+b,:]
        y = Y_test[r:r+b]
    return x,y
```

This function is used to retrieve batches of data for training or testing.

### 16.3.5 Model Definition

```
def model(x,w):
    return x@w[0]
```

This is a simple linear model, defined by matrix multiplication with the weight matrix ‘w’.

### 16.3.6 Plotting and Logging

```
def make_plots():
    acc_train = acc(model(x,w),y)
    xt,yt = get_batch('test')
    acc_test = acc(model(xt,w),yt)
    wb.log({"acc_train": acc_train, "acc_test":
           acc_test})
```

This function is used to make plots and log accuracy metrics to the wandb platform, useful for tracking model performance.

## 16.4 Recap

This chapter has provided a comprehensive guide to the code necessary for running the demo. The functions cover a wide range of tasks, from

data manipulation and visualization to defining and evaluating a simple linear model. The integration of the wandb platform further enables tracking and analysis of the model's performance.



## Chapter 17

# Training the Simple Linear Model on MNIST

## 17.1 Introduction

This chapter focuses on the actual training process of a simple linear model on the MNIST dataset. The code leverages the wandb platform for experiment tracking and utilizes the Adam optimizer for weight updates.

## 17.2 Setting Up the Experiment

### 17.2.1 Initializing Wandb

```
wb.init(project="Simple Linear MNIST ");  
c = wb.config
```

Here, the code initializes a new experiment in the wandb platform and configures it with the given project name. The configuration parameters are stored in the 'c' object.

### 17.2.2 Setting Hyperparameters

```
c.h = 0.01  
c.b = 1024  
c.epochs = 100000
```

This part sets the hyperparameters for the training process, including the learning rate ('c.h'), batch size ('c.b'), and the number of training epochs ('c.epochs').

## 17.3 Model Initialization and Optimization

### 17.3.1 Initializing Weights

```
w = [GPU(Truncated_Normal((784,10)))]
```

This code snippet initializes the weight matrix with a shape corresponding to the input size (784) and the number of classes (10), using the truncated normal distribution.

### 17.3.2 Setting Up the Optimizer

```
optimizer = torch.optim.Adam(w, lr=c.h)
```

The code leverages the Adam optimizer for training, which is known for its efficiency and adaptive learning rates.

## 17.4 Training Loop

```
for i in range(c.epochs):
    x,y = get_batch('train')
    loss = cross_entropy(softmax(model(x,w)),y)
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
    wb.log({"loss": loss})
    make_plots()
    if i % 1000 == 0 :
        montage_plot((w[0].T).reshape(10,28,28).cpu().
                     detach().numpy())
```

This code block represents the training loop where the following operations are performed:

1. Retrieving a batch of data using the 'get batch' function.
2. Calculating the loss using the cross-entropy between the softmax of the model's output and the target labels.
3. Zeroing out the gradients from previous iterations using 'optimizer.zero grad()'.

4. Computing the gradients with respect to the loss using `loss.backward()`.
5. Updating the weights using `optimizer.step()`.
6. Logging the loss to the wandb platform.
7. Making various plots including the ‘montage plot’ for the weights.

## 17.5 Recap

This chapter details the training process of a simple linear model on the MNIST dataset. Through a series of well-structured steps, including initialization, setting hyperparameters, defining the optimizer, and executing the training loop, this code offers a complete pipeline for training a linear classifier. The integration with the wandb platform enables efficient tracking of training progress and results. Visualization of the weights using a montage plot adds insights into the training process, offering a view into how the model’s weights evolve over time. It’s a robust example of how simple linear models can be applied to complex datasets like MNIST with the help of modern tools and libraries.





## Chapter 18

# Training a Fully Connected Neural Network on MNIST

### 18.1 Introduction

This chapter outlines the construction and training of a simple, fully connected neural network on the MNIST dataset. Two approaches are demonstrated: manual gradient descent and the utilization of the Adam optimizer.

### 18.2 Model Definition

#### 18.2.1 Gradient Step Function

```
def gradient_step(w):  
    for j in range(len(w)):  
        w[j].data = w[j].data - c.h*w[j].grad.data  
        w[j].grad.data.zero_()
```

This function performs a manual gradient descent step on the weights, updating the values and zeroing the gradients.

#### 18.2.2 ReLU Activation Function

```
def relu(x):  
    return x * (x > 0)
```

The ReLU activation function is defined, which applies the non-linearity element-wise to its input.

#### 18.2.3 Model Architectures

Three equivalent implementations of the model are given, showcasing different ways to structure the same computations:

```

def model(x,w):
    for j in range(len(w)):
        x = relu(x@w[j])
    return x

def model(x,w):
    return relu(relu(x@w[0])@w[1])

def model(x,w):
    x = x@w[0]
    x = relu(x)
    x = x@w[1]
    x = relu(x)
    return x

```

All three versions consist of two layers with ReLU activation, the difference being in how the operations are organized.

## 18.3 Training the Model

### 18.3.1 Training with Manual Gradient Descent

```

wb.init(project="Simple Fully Connected MNIST");
c = wb.config
c.h = 0.001
c.b = 100
c.layers = 2
c.epochs = 100000
c.f_n = [784,500,10]
w = [ GPU(randn_trunc((c.f_n[i],c.f_n[i+1]))) for i in
    range(c.layers) ]
for i in range(c.epochs):
    x,y = get_batch('train')
    loss = cross_entropy(softmax(model(x,w)),y)
    loss.backward()
    gradient_step(w)
    if (i+1) % 1 == 0:
        make_plots()

```

Here, the weights are initialized and the model is trained using the manually defined gradient descent. The training loop consists of computing the loss, performing backpropagation to compute the gradients, and updating the weights using the ‘gradient step’ function.

### 18.3.2 Training with the Adam Optimizer

```
wb.init(project="Simple Fully Connected MNIST");
c = wb.config
c.h = 0.001
c.b = 100
c.layers = 2
c.epochs = 100000
c.f_n = [784,500,10]
w = [ GPU(randn_trunc((c.f_n[i],c.f_n[i+1]))) for i in
      range(c.layers) ]
optimizer = torch.optim.Adam(w, lr=c.h)
for i in range(c.epochs):
    x,y = get_batch('train')
    loss = cross_entropy(softmax(model(x,w)),y)
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
    if (i+1) % 1 == 0:
        make_plots()
```

This segment of code is almost identical to the previous one but replaces the manual gradient descent step with the Adam optimizer. It continues to perform the same training operations but leverages the built-in Adam implementation in PyTorch for the weight updates.

## 18.4 Recap

This chapter illustrates the flexibility and choices available when constructing and training neural networks. By showcasing two different training approaches and multiple ways to define the same model architecture, it highlights the adaptability and diversity in building deep learning

models. The inclusion of both manual gradient descent and the popular Adam optimizer also demonstrates how these two methods can be applied interchangeably to the same problem, reflecting the rich landscape of optimization techniques in deep learning.

## Chapter 19

# Building an Autoencoder for Fashion MNIST

## 19.1 Introduction

This chapter presents the process of creating an autoencoder to reconstruct Fashion MNIST images. The autoencoder compresses the images into a lower-dimensional latent space and then decodes them back into the original space. Key parts of the code, including the definition of various functions and the main training loop, are detailed below.

## 19.2 Utility Functions

### 19.2.1 Truncated Normal Random Numbers

```
def randn_trunc(s):  
    mu = 0  
    sigma = 0.1  
    R = stats.truncnorm((-2*sigma - mu) / sigma, (2*  
        sigma - mu) / sigma, loc=mu, scale=sigma)  
    return R.rvs(s)
```

The function `randntrunc` creates random values from a truncated normal distribution, a normal distribution cut off at specified boundaries.

### 19.2.2 GPU Functions

```
def GPU(data):  
    return torch.tensor(data, requires_grad=True,  
        dtype=torch.float, device=torch.device('cuda'))
```

This function moves data to the GPU, converting it to a PyTorch tensor that requires gradients.

## 19.3 Data Loading and Preprocessing

Fashion MNIST data is loaded and preprocessed. The pixel values are normalized to lie within the range  $[-1, 1]$ , and only the images corresponding to a specific label ‘n’ are selected.

## 19.4 Building an Autoencoder

### 19.4.1 Model Definition

The autoencoder consists of two main parts: an encoder that compresses the images into a lower-dimensional latent space, and a decoder that reconstructs the original images.

```
def Encoder(x,w):
    return x@w[0]

def Decoder(x,w):
    return x@(w[0].T)

def Autoencoder(x,w):
    return Decoder(Encoder(x,w),w)
```

These functions take an input ‘x’ and a list of weights ‘w’ and apply a linear transformation.

### 19.4.2 Training the Autoencoder

The training loop iterates over the dataset for a specified number of steps. The images are passed through the autoencoder, and the mean squared error (MSE) between the reconstructed images and the original images is minimized.

```
w0 = GPU(np.random.randn(784,78))
w = [w0]
optimizer = torch.optim.Adam(params=w, lr=
    learning_rate, weight_decay=1e-5)
for step in range(steps):
    x,y = get_batch('train')
```

```
x2 = Autoencoder(x,w)
loss = MSE(x2, x)
optimizer.zero_grad()
loss.backward()
optimizer.step()
```

The Adam optimizer is used, and the training progress is optionally printed every 1000 steps.

### 19.4.3 Testing the Autoencoder

The code also includes a test section, where the trained autoencoder is applied to a batch of test images. The original, reconstructed, and difference images are visualized using the ‘montage plot’ function.

## 19.5 Recap

This chapter described the process of creating and training an autoencoder on the Fashion MNIST dataset. Through the use of several utility functions and a clear definition of the model and training loop, the code provides a complete example of how to compress and reconstruct images using neural networks. This simple yet powerful model serves as a foundational building block for more complex generative and unsupervised learning techniques.





## Chapter 20

# Convolutional Autoencoder (CAE) for Fashion MNIST

### 20.1 Introduction

The Convolutional Autoencoder (CAE) is a variant of the classic autoencoder that leverages convolutional layers. It's capable of capturing spatial patterns in data, which often results in better performance on image data like the Fashion MNIST dataset.

### 20.2 Reshaping and Normalizing Data

```
X = X.reshape(X.shape[0],1,28,28)
X_test = X_test.reshape(X_test.shape[0],1,28,28)
X = torchvision.transforms.functional.normalize(X
    ,0.5,0.5)
X_test = torchvision.transforms.functional.normalize(
    X_test,0.5,0.5)
```

The data is reshaped to fit the input format expected by convolutional layers and normalized.

### 20.3 Encoder Definition

```
def Encoder(x,w):
    x = relu(conv2d(x,w[0], stride=(2, 2), padding=(1,
        1)))
    x = relu(conv2d(x,w[1], stride=(2, 2), padding=(1,
        1)))
    x = x.view(x.size(0), 6272)
    x = linear(x,w[2])
    return x
```

The encoder consists of two convolutional layers followed by a linear layer. Strided convolutions are used to reduce the spatial dimensions, thereby compressing the data.

## 20.4 Decoder Definition

```
def Decoder(x,w):
    x = linear(x,w[3])
    x = x.view(x.size(0), 128, 7, 7)
    x = relu(conv_transpose2d(x,w[4], stride=(2, 2),
                             padding=(1, 1)))
    x = torch.tanh(conv_transpose2d(x,w[5], stride=(2,
        2), padding=(1, 1)))
    return x
```

The decoder part attempts to reconstruct the original data from the compressed form. It consists of one linear layer followed by two transposed convolutional layers. Transposed convolutions are used to increase the spatial dimensions, reconstructing the original shape of the data.

## 20.5 Training the Convolutional Autoencoder

```
for i in range(num_steps):
    x_real,y = get_batch('train')
    x_fake = Autoencoder(x_real,w)
    loss = torch.mean((x_fake - x_real)**2)
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
    if i % 100 == 0: print(loss.item())
```

The training loop is similar to the earlier example but adapted for the CAE. The loss function remains the mean squared error between the reconstructed and original images.

## 20.6 Evaluation and Visualization

After training, the model is tested on a batch of images, and the reconstruction error is calculated. The original and reconstructed images are plotted for visual inspection.

```
image_batch, y = get_batch('test')
image_batch_recon = Autoencoder(image_batch, w)
montage_plot(image_batch[0:25, 0, :, :].cpu().detach().
              numpy())
montage_plot(image_batch_recon[0:25, 0, :, :].cpu().
              detach().numpy())
```

## 20.7 Recap

This section introduced a Convolutional Autoencoder, demonstrating its ability to capture spatial patterns more effectively. By using convolutional layers in both encoding and decoding phases, the model was able to learn a more intricate representation of the Fashion MNIST data, yielding a better understanding of the inherent spatial structures in the images.



## Chapter 21

# Generative Adversarial Network (GAN) on Fashion MNIST

### 21.1 Introduction

Generative Adversarial Networks (GANs) consist of two neural networks, the generator (G) and the discriminator (D), trained simultaneously through a competitive process. The generator aims to create data that is similar to some real data, while the discriminator attempts to distinguish between real and generated data.

### 21.2 Preprocessing Data for Class Label $n$

```
n = 7
index = np.where(Y == n)
X = X[index]
index = np.where(Y_test == n)
X_test = X_test[index]
```

This code filters the dataset to include only images with the class label  $n$ , in this case, 7, and normalizes the pixel values.

### 21.3 Moving Data to GPU

```
X = GPU_data(X)
X_test = GPU_data(X_test)
Y = GPU_data(Y)
Y_test = GPU_data(Y_test)
```

This snippet moves the data onto the GPU, preparing it for accelerated computation.

## 21.4 Discriminator Definition

```
def D(x,w):
    x = relu(conv2d(x,w[0], stride=(2, 2), padding=(1,
        1)))
    x = relu(conv2d(x,w[1], stride=(2, 2), padding=(1,
        1)))
    x = x.view(x.size(0), 6272)
    x = linear(x,w[2])
    x = torch.sigmoid(x)
    return x
```

The discriminator takes in an image and outputs a scalar representing its probability of being real. It uses two convolutional layers followed by a linear layer.

## 21.5 Generator Definition

```
def G(x,w):
    x = linear(x,w[3])
    x = x.view(x.size(0), 128, 7, 7)
    x = relu(conv_transpose2d(x,w[4], stride=(2, 2),
        padding=(1, 1)))
    x = torch.tanh(conv_transpose2d(x,w[5], stride=(2,
        2), padding=(1, 1)))
    return x
```

The generator takes in a latent vector and transforms it into a synthetic image using a linear layer followed by two transposed convolutional layers.

## 21.6 Training the GAN

```
for i in range(steps):
    images,y = get_batch('train')
```

```
d_loss = binary_cross_entropy(D(images,w),
    real_labels) + binary_cross_entropy(D(G(z1[i],w),w),
    fake_labels)
d_optimizer.zero_grad()
d_loss.backward()
d_optimizer.step()

g_loss = binary_cross_entropy(D(G(z2[i],w),w),
    real_labels)
g_optimizer.zero_grad()
g_loss.backward()
g_optimizer.step()

if i % 200 == 0:
    out = G(z1[np.random.randint(steps)],w)
    montage_plot(out.view(batch_size,1,28,28).
        detach().cpu().numpy()[0:25,0,:,:])
```

The code snippet describes the adversarial training process, where the discriminator is trained to minimize the binary cross-entropy loss for correctly classifying real and generated images, while the generator is trained to fool the discriminator.

## 21.7 Visualizing Generated Images

```
noise = GPU_data(torch.randn(1,64))
output = G(noise,w)
plot(output[0,0])
```

This code generates a new image using the trained generator, allowing for visual inspection of the learned generative model.

## 21.8 Recap

This section introduced the implementation of a GAN specifically tailored for generating images of a particular class in the Fashion MNIST dataset. By leveraging convolutional layers in both the generator and

discriminator, the model has the capacity to capture the spatial patterns inherent to the images, leading to realistic image synthesis.



## Chapter 22

# Convolutional Neural Networks

## 22.1 Convolution

### 22.1.1 Introduction

Convolution is a mathematical operation that combines two signals to produce a third. In image processing, it's commonly used to detect specific patterns or features within an image, such as edges, textures, or specific objects.

### 22.1.2 Loading an Image

```
image = io.imread("http://ian-albert.com/games/  
    super_mario_bros_maps/mario-2-2.gif")  
image = image[:,0:700,:]  
plot(image)
```

Here, an image from the game Super Mario Bros is loaded. It's trimmed to the relevant part for the analysis.

### 22.1.3 Extracting the Coin Pattern

```
coin = image[185:200,224:239,:]  
plot(coin)
```

A specific coin pattern is extracted from the image to be used as a kernel for convolution.

### 22.1.4 Preprocessing the Image and Coin

```
image=np.mean(image,axis=2)  
coin=np.mean(coin,axis=2)  
image=scale1(image)
```

```
coin=scale1(coin)
image = image - np.mean(image)
coin = coin - np.mean(coin)
```

The image and coin pattern are converted to grayscale and scaled to a standard range. The mean is subtracted to center the data around zero.

### 22.1.5 Performing Convolution

```
from scipy import signal
z=signal.convolve2d(image, np.rot90(coin, 2))
plot(z)
```

2D convolution is performed using the coin pattern as a kernel. The result is a new image where high intensity corresponds to the areas where the coin pattern is detected.

### 22.1.6 Visualizing the Coin Detection

```
plot(z==np.max(z))
[y,x] = np.where(z == np.amax(z))
plt.plot(x,-y,'.')
fig, ax = plt.subplots()
im = ax.imshow(image, cmap = 'gray')
ax.axis('off')
fig.set_size_inches(10, 10)
ax.scatter(x-6, y-6, c='#f97306', s=40)
plt.show()
```

This code visualizes the location of the coin pattern within the original image. It does this by finding the maximum value in the convolution result and plotting that location on the original image.

### 22.1.7 Recap

This section introduced convolution, a fundamental operation in image processing, and demonstrated how it can be used to detect specific patterns within an image. Convolution is not only a powerful tool for image

analysis but also the building block of convolutional neural networks (CNNs), which have become vital in various computer vision applications.

## 22.2 Applying a Convolutional Filter

### 22.2.1 Introduction

In this section, we demonstrate how to apply a convolutional filter to an image. Convolutional filters are fundamental in image processing and computer vision, and they play a critical role in extracting features and patterns from images.

### 22.2.2 Loading and Preprocessing the Image

```
image = io.imread("http://harborparkgarage.com/img/
    venues/pix/aquarium2.jpg")
image = image[:,:,:]
plot(image)
image = image.astype(float)/255.0
image.shape
```

We begin by loading an image and normalizing the pixel values to the range  $[0, 1]$ . This prepares the image for further processing.

### 22.2.3 Creating Random Filters

```
plot(np.random.random((5,5,3)))
filters = np.random.random((96,11,11,3))
plot(filters[1,:,:,:])
f = np.random.random((11,11,3))
plot(f)
```

We create a set of random filters for demonstration purposes. In practice, these filters would typically be learned by a convolutional neural network (CNN) from training data.

### 22.2.4 Transposing the Image

```
image = np.transpose(image, (2, 0, 1))
```

The image is transposed to have the channels as the first dimension, conforming to the input format expected by PyTorch’s convolutional functions.

### 22.2.5 Applying the Convolutional Filter

```
import torch.nn.functional as F
import torch

f = np.random.random((1,3,11,11))
image = image[None,:,:,:]
f = torch.from_numpy(f)
image = torch.from_numpy(image)
image2 = F.conv2d(image,f)
image2.shape
image2 = image2.numpy()
plot(image2[0,0,:,:])
```

We apply the convolutional filter using PyTorch’s functional API. First, the filter is reshaped to have the proper dimensions (batch size, number of channels, height, width), and both the image and filter are converted to PyTorch tensors. Then, the ‘F.conv2d’ function is used to apply the convolution, and the result is plotted.

### 22.2.6 Recap

This section provided a hands-on demonstration of applying a convolutional filter to an image using PyTorch. Convolutional filters are a cornerstone of modern computer vision and the foundation of CNNs. By understanding how to manually apply these filters, we gain insight into the operations that CNNs perform as they process images.

## 22.3 Applying Various Convolution Filters

### 22.3.1 Introduction

Convolution filters are used to perform a wide variety of operations in image processing, such as edge detection, noise reduction, and feature extraction. This section demonstrates how different types of filters affect an image.

### 22.3.2 Loading and Preprocessing the Image

```
image = io.imread("https://tinyurl.com/SpaceShip1234")
image = image[:, :, :]
plot(image)
image.shape
image = np.mean(image, axis=2)
plot(image)
```

The image is loaded, and the mean of the RGB channels is computed to create a grayscale image.

### 22.3.3 Applying Handcrafted Filters

```
a = np.matrix([[1,2,1],[0,0,0],[-1,-2,-1]])
y = signal.convolve2d(image, a, mode='same')
plot(y)
a = np.transpose(a)
y = signal.convolve2d(image, a, mode='same')
plot(y)
```

Here, a Sobel filter is applied to the image in both horizontal and vertical orientations to detect edges. The filter is then transposed to demonstrate the effect on the image.

### 22.3.4 Applying Random Filters

```
b = np.random.random((25,25))
y = signal.convolve2d(image, b)
plot(y)
```

A random filter is applied to the image, demonstrating how different and unpredictable effects can be achieved.

### 22.3.5 Implementing a Custom Convolution Function

```
def conv2(x,f):
    x2 = np.zeros(x.shape)
    for i in range(1,x.shape[0]-1):
        for j in range(1,x.shape[1]-1):

            x2[i,j] = f[0,0] * x[i-1,j-1] \
+                    f[0,1] * x[i-1,j]    \
+                    f[0,2] * x[i-1,j+1]  \
+                    f[1,0] * x[i,j-1]    \
+                    f[1,1] * x[i,j]      \
+                    f[1,2] * x[i,j+1]    \
+                    f[2,0] * x[i+1,j-1]  \
+                    f[2,1] * x[i+1,j]    \
+                    f[2,2] * x[i+1,j+1]

    return x2

a = 5 * np.random.random((3,3)) - 5 * np.random.random
    ((3,3))
z = conv2(x,a)
plot(x)
plot(z)
```

A custom convolution function is defined and applied to an image with a randomly generated filter. This helps demonstrate the manual computation of a convolution.

### 22.3.6 Experimenting with Various Random Filters

## 22.4. COMPARING PERFORMANCE OF DIFFERENT CONVOLUTION METHODS

```
for i in range(9):
    a = 2*np.random.random((3,3))-1
    print(a)
    z = conv2(x,a)
    plot(z)
```

Finally, the code iterates through a series of randomly generated filters, applying each one to the image and plotting the result. This showcases the wide range of effects that can be achieved with different filters.

### 22.3.7 Recap

Through this hands-on demonstration, we have explored the application of various convolutional filters to images, including handcrafted, random, and custom-defined filters. By visualizing the effects of these filters, we gain an understanding of how convolutional filters work in image processing and deepen our insights into their role in feature extraction and pattern recognition in machine learning models.

## 22.4 Comparing Performance of Different Convolution Methods

### 22.4.1 Introduction

The efficiency of convolution operations is vital in image processing and deep learning, especially when dealing with large-scale data. This section will compare the time taken to perform convolutions using a custom-defined function, SciPy's built-in function, and PyTorch's GPU-accelerated method.

### 22.4.2 Using a Custom-Defined Function

```
a = 2*np.random.random((9,3,3))-1
start_time = time.time()
for i in range(9):
    z=conv2(x,a[i,:,:])
```

```
print("--- %s seconds ---" % (time.time() - start_time))
```

Here, the custom-defined function `conv2` is used to apply a series of 9 randomly generated filters to an image. The time taken to execute this operation is printed.

### 22.4.3 Using SciPy's Convolution Function

```
a = 2*np.random.random((9,3,3))-1
start_time = time.time()
for i in range(9):
    z = signal.convolve2d(x,a[i,:,:])
print("--- %s seconds ---" % (time.time() - start_time))
```

In this segment, the same operation is performed using SciPy's `convolve2d` function. The execution time is again measured and printed, enabling a direct comparison with the custom-defined function.

### 22.4.4 Using GPU Acceleration with PyTorch

```
#GPU Processing Timing, No Loop, 96 filters!!
a2 = 2*np.random.random((96,1,3,3))-1
x2 = torch.tensor(x).cuda()
a2 = torch.tensor(a2).cuda()
x2 = x2[None,None,:,:]
start_time = time.time()
z = torch.nn.functional.conv2d(x2,a2)
print("--- %s seconds ---" % (time.time() - start_time))
```

This segment showcases the efficiency of GPU acceleration by applying 96 randomly generated filters to the image using PyTorch's `conv2d` function on the GPU. The time taken is significantly reduced compared to the other methods.



### 22.4.5 Recap

By comparing these three methods, it's evident that the efficiency of convolution operations can vary significantly based on the implementation and resources used. The use of specialized libraries and hardware acceleration (such as GPU processing) can drastically reduce computation time, thereby enhancing the feasibility of large-scale image processing and deep learning tasks. It emphasizes the importance of selecting the right tools and methods for specific tasks to achieve optimal performance.



## Chapter 23

# Image Classification Using Pre-trained AlexNet

### 23.0.1 Introduction

AlexNet is a well-known deep convolutional neural network that has been instrumental in advancing the field of deep learning. With the PyTorch library, developers can quickly utilize this pre-trained model to classify images. The following sections describe how to preprocess an image and use AlexNet to predict its class.

### 23.0.2 Defining the AlexNet Model

```
import torch
from torchvision import models

#define alexnet model
alexnet = models.alexnet(pretrained=True).cuda(0)
```

Here, the code imports the AlexNet model from the PyTorch's torchvision package and downloads the pre-trained weights. The `.cuda(0)` method ensures that the model will run on the GPU, if available, for faster computation.

### 23.0.3 Loading the Class Labels

```
labels = {int(key):value for (key, value) in requests.
    get('https://s3.amazonaws.com/mlpipes/pytorch-quick
    -start/labels.json').json().items() }
```

The class labels corresponding to the output indices of the model are loaded from a JSON file. These labels will be used to interpret the prediction made by the model.

### 23.0.4 Image Preprocessing

```
from torchvision import transforms
from PIL import Image

#transform image for use in model
preprocess = transforms.Compose([
    transforms.Resize(256),
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std
        =[0.229, 0.224, 0.225])
])
```

Before using the image with AlexNet, it must be preprocessed. The code defines a sequence of transformations to resize, crop, and normalize the image, following the exact preprocessing steps that were used when training AlexNet.

### 23.0.5 Loading and Classifying an Image

```
#load the image from its url
url = 'https://tinyurl.com/DogImage1234'
img = Image.open(requests.get(url, stream=True).raw)
img_t = preprocess(img).unsqueeze_(0).cuda(0)

#classify the image with alexnet
scores, class_idx = alexnet(img_t).max(1)
print('Predicted class:', labels[class_idx.item()])
```

The image is loaded from a URL, preprocessed, and passed through the AlexNet model. The resulting prediction is matched with the corresponding class label, and the predicted class is printed.

### 23.0.6 Recap

This example demonstrates how to leverage a pre-trained deep learning model like AlexNet in PyTorch to perform image classification. With

just a few lines of code, it becomes possible to utilize the power of deep learning to recognize complex patterns in images, showing the flexibility and power of modern deep learning frameworks.

## 23.1 Exploring and Manipulating AlexNet Layers

### 23.1.1 Introduction

AlexNet consists of several layers of convolutions, activations, pooling, and fully connected operations. Understanding how data is transformed through these layers is essential for interpreting deep neural networks. This section walks through the layers of AlexNet, applying them manually to an image, and explores how to save and load model weights.

### 23.1.2 Extracting Weights from AlexNet

```
w0 = alexnet.features[0].weight.data
w1 = alexnet.features[3].weight.data
w2 = alexnet.features[6].weight.data
w3 = alexnet.features[8].weight.data
w4 = alexnet.features[10].weight.data
w5 = alexnet.classifier[1].weight.data
w6 = alexnet.classifier[4].weight.data
w7 = alexnet.classifier[6].weight.data

w = [w0, w1, w2, w3, w4, w5, w6, w7]

torch.save(w, 'Hahn_Alex.pt')
```

Weights of individual layers are extracted and collected into a list. They are then saved to a file for future reference. This is useful for preserving the exact state of a model or for sharing model weights.

### 23.1.3 Loading Weights

```
w = torch.load('Hahn_Alex.pt')
[w0,w1,w2,w3,w4,w5,w6,w7] = w
```

The saved weights can be loaded back into Python for use. This allows the state of the model to be restored or transferred between different systems or sessions.

### 23.1.4 Applying Layers Manually

The following code applies each layer of AlexNet manually to the image, using the extracted weights:

```
import torch.nn.functional as F

f0 = F.conv2d(img_t, w0, stride=4, padding=2)

f1 = F.relu(f0)

f2 = F.max_pool2d(f1, kernel_size=3, stride=2, padding
                  =0, dilation=1)

f3 = F.conv2d(f2, w1, stride=1, padding=2)

f4 = F.relu(f3)

f5 = F.max_pool2d(f4, kernel_size=3, stride=2, padding
                  =0, dilation=1)

f6 = F.conv2d(f5, w2, stride=1, padding=1)

f7 = F.relu(f6)

f8 = F.conv2d(f7, w3, stride=1, padding=1)

f9 = F.relu(f8)

f10 = F.conv2d(f9, w4, stride=1, padding=1)

f11 = F.relu(f10)
```

```
f12 = F.max_pool2d(f11, kernel_size=3, stride=2,
                    padding=0, dilation=1)

f13 = F.adaptive_avg_pool2d(f12, output_size=6).flatten()

f14 = F.linear(f13, w5)

f15 = F.relu(f14)

f16 = F.linear(f15, w6)

f17 = F.relu(f16)

f18 = F.linear(f17, w7)

out = f18.argmax().item()

labels[out]
```

This step-by-step application of each layer provides a detailed look at how the input image is transformed through the network. This can be valuable for educational purposes or for in-depth analysis of the model's behavior.

### 23.1.5 Recap

This section showed how to extract, save, and load weights from the AlexNet model, and how to manually apply each layer to an input image. This understanding of the inner workings of a neural network is vital for anyone looking to modify, analyze, or understand these powerful models. It also highlights the flexibility and ease with which one can manipulate complex deep learning models using PyTorch.

## 23.2 Understanding ReLU and Max Pooling

### 23.2.1 ReLU: Rectified Linear Unit

#### Definition

ReLU is a non-linear activation function that introduces non-linearity into the network. It's defined as:

$$f(x) = \max(0, x)$$

In other words, ReLU replaces all negative values in a tensor with zero.

#### Usage in CNN

ReLU is used to add the capability to model complex, non-linear relationships between inputs and outputs. It's typically applied after each convolution operation.

```
f1 = F.relu(f0)
```

Here, 'f0' is the output from a convolution layer, and 'f1' is the result after applying the ReLU activation function.

#### Advantages

ReLU's main advantages are computational efficiency and the ability to mitigate the vanishing gradient problem, a common issue in training deep networks.

### 23.2.2 Max Pooling

#### Definition

Max Pooling is a down-sampling operation that selects the maximum value from a group of values. In CNNs, it's usually applied to small regions of the input feature map.



### Usage in CNN

Max Pooling is used to reduce the spatial dimensions of the feature maps, thereby reducing the number of parameters and computational cost. It's typically applied after one or more convolution and activation layers.

```
f2 = F.max_pool2d(f1, kernel_size=3, stride=2, padding=0, dilation=1)
```

Here, 'f1' is the output from a ReLU layer, and 'f2' is the result after applying Max Pooling with a 3x3 kernel and a stride of 2.

### Advantages

Max Pooling has several advantages, including:

- Reducing the number of parameters, which helps prevent overfitting.
- Preserving the most significant information, as it keeps only the maximum values.
- Providing a form of translation invariance, as the exact position of the features within the pooling region becomes less important.

### 23.2.3 Recap

ReLU and Max Pooling are fundamental operations in CNN architectures. ReLU introduces non-linearity, allowing the model to learn more complex patterns, while Max Pooling helps to reduce dimensions and computational cost. Together, they play a vital role in making deep learning models both powerful and efficient. Understanding these operations is essential for anyone working with CNNs, as it provides insight into how these models process and learn from input data.

## 23.3 Training AlexNet with Custom Dataset

### 23.3.1 Mounting Google Drive

The code begins by mounting Google Drive to access the dataset located in the Google Drive folder. This step is necessary when running the code in a Google Colab environment.

```
drive.mount('/content/gdrive')
```

### 23.3.2 Loading and Pre-processing Data

The images are loaded from specific directories and pre-processed using transformations like resizing, cropping, flipping, and normalizing.

```
data_transforms = {
    'train': transforms.Compose([
        transforms.RandomResizedCrop(224),
        transforms.RandomHorizontalFlip(),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406],
                              [0.229, 0.224, 0.225])
    ]),
    'valid': transforms.Compose([
        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406],
                              [0.229, 0.224, 0.225])
    ]),
}

image_datasets = {x: datasets.ImageFolder(data_dir + '/' + x, data_transforms[x]) for x in ['train', 'valid']}

dataloaders = {x: DataLoader(image_datasets[x],
                              batch_size=16, shuffle=True, num_workers=4) for x in ['train', 'valid']}
```

### 23.3.3 AlexNet Architecture

The pre-trained AlexNet architecture is loaded, and the final layer is customized to match the number of classes in the dataset. The model is also transferred to the GPU if available.

```
alexnet = models.alexnet(pretrained=True)
num_ftrs = alexnet.classifier[6].in_features
alexnet.classifier[6] = nn.Linear(num_ftrs, len(
    class_names))
alexnet = alexnet.to(device)
```

### 23.3.4 Training Process

The model is trained for a specified number of epochs using Stochastic Gradient Descent (SGD) as the optimizer and CrossEntropyLoss as the loss function. The training loop involves both training and validation phases, and the best parameters are saved.

```
num_epochs = 10
acc_best = 0
parameters_best = 0

for epoch in range(num_epochs):
    print(epoch, " of ", num_epochs - 1)
    print('-' * 10)

    running_corrects = 0

    for inputs, labels in dataloaders["train"]:
        inputs = inputs.to(device)
        labels = labels.to(device)

        outputs = alexnet(inputs)
        preds = torch.max(outputs, 1)[1]

        optimizer.zero_grad()
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
```

```

        running_corrects += torch.sum(preds == labels.
            data)

    print('Train Acc: {:.4f}'.format(running_corrects
        / dataset_sizes["train"]))

    # Evaluate the AlexNet model on Validation Data
    alexnet.eval()

    running_corrects = 0

    for inputs, labels in dataloaders["valid"]:
        inputs = inputs.to(device)
        labels = labels.to(device)

        outputs = alexnet(inputs)
        preds = torch.max(outputs, 1)[1]

        running_corrects += torch.sum(preds == labels.
            data)

    acc_valid = running_corrects / dataset_sizes["
        valid"]

    if acc_valid > acc_best:
        acc_best = acc_valid
        parameters_best = alexnet.parameters()

    print('Valid Acc: {:.4f}'.format(acc_valid))
    if acc_valid > 0.99:
        print("Done!")
        break

alexnet.parameters = parameters_best
print('Training complete')
```

### 23.3.5 Analysis and Summary

This code snippet showcases the complete process of fine-tuning a pre-trained AlexNet model for a custom image classification task. By utilizing a well-known architecture like AlexNet and adapting it to a specific dataset, one can achieve robust classification performance without having to design a model from scratch.

The code leverages powerful PyTorch functionality, including data transformations, pre-trained models, loss functions, and optimization techniques. The inclusion of validation and best-model saving strategies further aligns this example with best practices in deep learning model training.

The fine-tuning of pre-trained models like AlexNet represents an important approach in transfer learning, enabling the utilization of pre-existing knowledge to new, related tasks, thus often saving time and computational resources.

## 23.4 Model Saving, Loading, and Inference

### 23.4.1 Saving the Trained Model

After training the AlexNet model, it is saved to a specified path in Google Drive. This will enable the model to be reloaded and used for later predictions without needing to retrain.

```
model_path = '/content/gdrive/MyDrive/Data2023/
              ballgame/alexnet_classification_ballgame.pth'
torch.save(alexnet.state_dict(), model_path)
```

### 23.4.2 Loading the Trained Model

A function, `load_model`, is defined to load the saved model from the specified path. It recreates the AlexNet architecture, loads the saved parameters, and sets the model in evaluation mode.

```
def load_model(model_path):
    model = models.alexnet()
```

```

num_ftrs = model.classifier[6].in_features
model.classifier[6] = nn.Linear(num_ftrs, len(
    class_names))
model.load_state_dict(torch.load(model_path))
model.eval()
return model.to(device)

```

### 23.4.3 Preprocessing a New Image

A function, `preprocess_image`, is created to load an image from a URL and preprocess it using the same transformations that were applied to the training data.

```

def preprocess_image(url, transform):
    response = requests.get(url)
    img = Image.open(BytesIO(response.content)).
        convert("RGB")
    img_tensor = transform(img)
    return img_tensor.unsqueeze(0).to(device)

```

### 23.4.4 Performing Inference

The `predict_image_url` function takes a URL of an image and the loaded model, preprocesses the image, and runs it through the model to get a prediction.

```

def predict_image_url(url, model):
    img_tensor = preprocess_image(url, data_transforms
        ['valid'])
    output = model(img_tensor)
    pred = torch.max(output, 1)[1]
    return class_names[pred]

image_url = 'https://tinyurl.com/DogImage1234'
prediction = predict_image_url(image_url,
    trained_model)
print("The predicted class for the input image is:",
    prediction)

```

### 23.4.5 Analysis and Summary

This section of the code illustrates how to save a trained deep learning model and then leverage it for making predictions on unseen data. Saving and loading models are crucial steps in deploying machine learning models, allowing trained models to be shared, evaluated, or put into production.

The code also highlights the importance of preprocessing new input data in the same manner as the training data. This consistency ensures that the model receives the correct input format and can make meaningful predictions.

Lastly, this example underscores the practicality and applicability of deep learning models, where a trained model can readily make predictions on real-world data such as an image fetched from a URL. Such a flow is typical in real-world applications where models are trained, saved, and later used to make predictions.





# Chapter 24

## Recap

### 24.1 Reflection on Key Concepts

Throughout this book, we have explored the fascinating world of deep learning, specifically focusing on PyTorch, one of the most powerful and widely-used deep learning frameworks. From the fundamentals of tensors and gradient-based optimization to advanced architectures like AlexNet, we have delved into the underlying principles and hands-on applications that make deep learning a transformative technology.

We covered various concepts such as:

- Basic tensor operations and the autograd system for automatic differentiation.
- Building custom neural network architectures and training them on real datasets.
- Transfer learning and the utilization of pre-trained models to enhance performance.
- Image preprocessing, transformation, and the application of convolutional layers.
- Saving, loading, and deploying trained models for real-world inference.

### 24.2 Future Perspectives

The landscape of deep learning is vast, and what has been covered in this book is just the tip of the iceberg. The field continues to advance at a rapid pace, with new research, models, and applications emerging continually. Readers are encouraged to:

- Explore specialized domains like natural language processing, reinforcement learning, and generative models.
- Engage with the community by contributing to open-source projects and participating in hackathons and challenges.
- Stay up-to-date with the latest research, techniques, and best practices through conferences, workshops, and online resources.
- Consider ethical implications and ensure that the developed models are fair, transparent, and responsible.

## 24.3 Final Thoughts

This book has aimed to be a comprehensive guide for both beginners and intermediate practitioners in deep learning. By bridging theory with practical examples, we hope to have provided the reader with a solid foundation and the inspiration to explore further.

Deep learning is more than just a technological trend; it's a tool with the potential to redefine how we understand, interact with, and shape our world. The journey doesn't end here; it's just the beginning.

Whether you're pursuing a career in artificial intelligence, embarking on a research project, or simply indulging in personal curiosity, we hope this book serves as a valuable companion on your journey through the dynamic and exciting world of deep learning.

Thank you for embarking on this learning adventure with us.

**Happy Coding!**